

Macro forecasting: ARDL-MIDAS regressions

title: Macro forecasting: ARDL-MIDAS regressions autosize: true

ARDL-MIDAS regression

ARDL-MIDAS(p_y^Q, q_X^M) model

$$Y_{t+h}^{Q,h} = \mu_h + \sum_{j=0}^{p_y^Q-1} \rho_{j+1}^h Y_{t-j}^Q + \beta^h \sum_{j=0}^{q_X^M-1} \sum_{i=0}^{m-1} \omega_{i+j*m}(\theta_h) X_{m-i,t-j}^M + \epsilon_{t+h}^h \quad (1)$$

where

- $Y_{t+h}^{Q,h}$ — low-frequency (quarterly) target variable
- $X_{m-i,t-j}^M$ — high-frequency (monthly) predictor variable
- $\omega_{i+j*m}(\theta_h)$ — weight function
- $\Theta := (\mu_h, \rho_1, \dots, \rho_{p_y^Q}, \beta^h, \theta_h)$ — model parameters estimated by NLS

Macro forecasting

Preliminaries

```
# Clean workspace
rm(list=ls())
# Load required package
require("MIDASLec")
# Load data for macro examples:
data("example1")
```

Check if data is loaded in your RStudio (Environment)

You should see: ads, cfnai, payems, rgdp

Trasform data to growth rates

ADS and CFNAI are macro factors, we don't need to transform them

payems:

```
payems[-1, 2] <- log(payems[-1, 2]/payems[-dim(payems)[1], 2])*100
payems <- payems[-1, ]
```

you can plot it to check if it looks as expected

Trasform data to growth rates

```
plot(payems[,1], payems[,2], type='l', ylab='Monthly Nonfarm Payrolls growth rate', xlab='Months')
```

Trasform data to growth rates

rgdp:

```
rgdp[-1, 2] <- log(rgdp[-1, 2]/rgdp[-dim(rgdp)[1], 2])*100
rgdp <- rgdp[-1, ]
```

MIDAS in R

There are several packages in R to run different MIDAS regressions.

For this example we will rely on midasr package with some additional functions in MIDASLec.

Nonfarm Payrolls, construct MIDAS data structures which will be used in MIDAS models.

First example is with Nonfarm Payrolls data

Set initial and last date for in-sample estimation:

```
est.start <- "1985-01-01"
est.end <- "2009-01-01"
```

Nonfarm Payrolls example

Get data structures needed for MIDAS regression:

```
data.payems.in <- mixed.freq.data(rgdp[,2], rgdp[,1],
                                payems[,2], payems[,1], x.lag=9,
                                y.lag=1, horizon=3, est.start=as.Date(est.start),
                                est.end=as.Date(est.end), disp.flag = TRUE)
```

Frequency of Data Y: 3 month(s)

Frequency of Data X: 1 month(s)

Start Date: 1985-01-01

Terminal Date: 2009-01-01

Mixed frequency regression time frame:

Reg Y(1985-03-01)`s on: Y(1984-12-01)`s X(1984-12-01)`s X(1984-11-01)`s ... X(1984-04-01)`s

Reg Y(1985-06-01)`s on: Y(1985-03-01)`s X(1985-03-01)`s X(1985-02-01)`s ... X(1984-07-01)`s

Reg Y(2009-03-01)`s on: Y(2008-12-01)`s X(2008-12-01)`s X(2008-11-01)`s ... X(2008-04-01)`s

For this example we use exponential Almon weights

```
weight <- nealmon
```

Get a set of good starting values for nonlinear optimization

```
set.seed(123)
startx.all <- get.start.adl.midas(y=data.payems.in$est.y,
                                X=data.payems.in$est.x, z=data.payems.in$est.lag.y,
                                weight=weight, par.num.weight=3, num.evals=10000, num.coef=1)
```

Estimate ARDL-MIDAS regression using *midas_r_plain* function from midasr package

```
payems.est.midas <- midas_r_plain(y=data.payems.in$est.y, X=data.payems.in$est.x,
                                z=cbind(data.payems.in$est.lag.y,
                                rep(1,times=length(data.payems.in$est.y))), weight=weight,
                                startx=startx.all[-c(4,5)], startz=startx.all[c(4,5)]
                                ,control=list(maxit=500))
```

Estimate ARDL regression model using time-averaged Nonfarm Payrolls data

```
payems.est.ta <- lm(data.payems.in$est.y ~ data.payems.in$est.lag.y +
                    rowMeans(data.payems.in$est.x[,1:3]) +
                    rowMeans(data.payems.in$est.x[,4:6]) +
                    rowMeans(data.payems.in$est.x[,7:9]))
```

Compare in-sample performance

```
sqrt(mean(payems.est.midas$residuals^2))
```

```
[1] 0.5103188
```

```
sqrt(mean(payems.est.ta$residuals^2))
```

```
[1] 0.4945295
```

Compare out-of-sample performance

```
payems.midas.oos <- forecast.adl(payems.est.midas, weight=weight,
                                par.num.weight=3, data.payems.in, is.intercept=TRUE)
payems.ta.oos <- forecast.ta.example(payems.est.ta,
                                    data.payems.in)
```

```
payems.midas.oos$rmse
```

```
[1] 0.4461215
```

```
payems.ta.oos$rmse
```

```
[1] 0.4486757
```

Plot lag polynomials

```
par(mfrow=c(1,2))
```

```
plot( weight(payems.est.midas$coefficients[c(1,2,3)], d = 9), type = 'l',
      xlab='Lag', ylab='Coefficient',
      main='Normalized exponential Almon')
```

```
plot( c(rep(as.numeric(payems.est.ta$coefficients[3]), times = 3),
        rep(as.numeric(payems.est.ta$coefficients[4]), times = 3),
        rep(as.numeric(payems.est.ta$coefficients[5]), times = 3)), type = 'l',
      xlab='Lag', ylab='Coefficient', main='Time-averaged data coefficients')
```

We will compare fixed, rolling, expanding windows predictions of GDP using these two models.

Fixed window

MIDAS

```
payems.midas.obj.fixed <- midas.adl(data.y = rgdp[,2], data.ydate = rgdp[,1],
                                   data.x = payems[,2], data.xdate = payems[,1],
                                   est.start = as.Date(est.start), est.end = as.Date(est.end),
                                   horizon = 3, x.lag = 9, y.lag = 1,
```

```
polynomial = "nealmon", method = "fixed",
disp.flag = FALSE, num.evals = 10000, num.coef = 1)
```

Fixed window

Time-averaged data

```
payems.ta.obj.fixed <- midas.adl(data.y = rgdp[,2], data.ydate = rgdp[,1],
data.x = payems[,2], data.xdate = payems[,1],
est.start = as.Date(est.start), est.end = as.Date(est.end),
horizon = 3, x.lag = 9, y.lag = 1,
polynomial = "timeaverage", method = "fixed",
disp.flag = FALSE)
```

We compared predictions already for fixed scheme

Rolling window

MIDAS

```
payems.midas.obj.rolling <- midas.adl(data.y = rgdp[,2], data.ydate = rgdp[,1],
data.x = payems[,2], data.xdate = payems[,1],
est.start = as.Date(est.start),
est.end = as.Date(est.end), horizon = 3,
x.lag = 9, y.lag = 1,
polynomial = "nealmon", method = "rolling",
disp.flag = FALSE, num.evals = 10000,
num.coef = 1)
```

Time-averaged data

```
payems.ta.obj.rolling <- midas.adl(data.y = rgdp[,2], data.ydate = rgdp[,1],
data.x = payems[,2], data.xdate = payems[,1],
est.start = as.Date(est.start),
est.end = as.Date(est.end), horizon = 3,
x.lag = 9, y.lag = 1, polynomial = "timeaverage",
method = "rolling",
disp.flag = FALSE)
```

Compare predictions for rolling window

```
payems.midas.obj.rolling$pred.obj$rmse
```

```
[1] 0.4284624
```

```
payems.ta.obj.rolling$pred.obj$rmse
```

```
[1] 0.4325905
```

Expanding window

MIDAS

```
payems.midas.obj.expand <- midas.adl(data.y = rgdp[,2], data.ydate = rgdp[,1],
  data.x = payems[,2], data.xdate = payems[,1],
  est.start = as.Date(est.start),
  est.end = as.Date(est.end), horizon = 3, x.lag = 9,
  y.lag = 1, polynomial = "nealmon",
  method = "expand", disp.flag = FALSE,
  num.evals = 10000, num.coef = 1)
```

Time-averaged data

```
payems.ta.obj.expand <- midas.adl(data.y = rgdp[,2], data.ydate = rgdp[,1],
  data.x = payems[,2], data.xdate = payems[,1],
  est.start = as.Date(est.start),
  est.end = as.Date(est.end), horizon = 3, x.lag = 9,
  y.lag = 1, polynomial = "timeaverage",
  method = "expand", disp.flag = FALSE)
```

Compare predictions for expanding window

```
payems.midas.obj.expand$pred.obj$rmse
```

```
[1] 0.4295506
```

```
payems.ta.obj.expand$pred.obj$rmse
```

```
[1] 0.4369391
```

CFNAI example

Set weighting function, sort data and get good initial values

```
weight <- nealmon
est.start <- "1987-01-01"
est.end <- "2011-12-01"
```

```
data.cfani.in <- mixed.freq.data(rgdp[,2], rgdp[,1],
  cfnaï[,2], cfnaï[,1], x.lag=12,
  y.lag=1, horizon=3, est.start=as.Date(est.start),
  est.end=as.Date(est.end),
  disp.flag = FALSE)
```

```
set.seed(123)
```

```
startx.all <- get.start.adl.midas(y=data.cfani.in$est.y, X=data.cfani.in$est.x,
  z=data.cfani.in$est.lag.y, weight=weight,
  par.num.weight=3, num.evals=10000, num.coef=1)
```

Estimate MIDAS and U-MIDAS regressions

```
cfnaï.est.midas <- midas_r_plain(y = data.cfani.in$est.y, X = data.cfani.in$est.x,
  z = cbind(data.cfani.in$est.lag.y,
  rep(1,times = length(data.cfani.in$est.y))),
```

```

weight=weight, startx = startx.all[-c(4,5)],
startz = startx.all[c(4,5)],
control = list(maxit = 1000))

cfnai.est.umidas <- lm(data.cfani.in$est.y ~ data.cfani.in$est.lag.y + data.cfani.in$est.x)
Compare in-sample performance
sqrt(mean(cfnai.est.midas$residuals^2))
[1] 0.4848867
sqrt(mean(cfnai.est.umidas$residuals^2))
[1] 0.4551793
Compare out-of-sample performance
cfnai.midas.oos <- forecast.adl(cfnai.est.midas,
weight=weight, par.num.weight=3,
data.cfani.in, is.intercept=TRUE)

cfnai.umidas.oos <- forecast.umidas(cfnai.est.umidas, data.cfani.in)
cfnai.midas.oos$rmse
[1] 0.3154912
cfnai.umidas.oos$rmse
[1] 0.3367636
Plot lag polynomials: exponential Almon
plot(weight(cfnai.est.midas$coefficients[c(2,3)], d = 12),
type = 'l', xlab = 'Lag', ylab = 'Coefficient',
main = 'Normalized exponential Almon lag polynomial')
Plot lag polynomials: unrestricted
plot(coef(cfnai.est.umidas),
type = 'l', xlab = 'Lag', ylab = 'Coefficient',
main = 'U-MIDAS lag polynomial')

```

ADS example

Set sample dates

```

est.start <- "1987-01-01"
est.end <- "2011-12-01"

```

Compute MIDAS forecasts

```

ads.midas.obj.fixed.for <- midas.adl(data.y = rgdp[,2], data.ydate = rgdp[,1],
data.x = ads[,2], data.xdate = ads[,1],
est.start = as.Date(est.start),
est.end = as.Date(est.end), horizon = 66, x.lag = 66,
y.lag = 1, polynomial = "nealmon",
method = "fixed", disp.flag = TRUE,
num.evals = 10000, num.coef = 10)

```

Compute U-MIDAS forecasts

```
ads.umidas.obj.fixed.for <- midas.obj.fixed.for <- midas.adl(data.y = rgdp[,2],
  data.ydate = rgdp[,1], data.x = ads[,2], data.xdate = ads[,1],
  est.start = as.Date(est.start),
  est.end = as.Date(est.end), horizon = 66, x.lag = 66,
  y.lag = 1, polynomial = "umidas",
  method = "fixed", disp.flag = TRUE)
```

Compute MIDAS nowcasts

```
ads.midas.obj.fixed.now <- midas.adl(data.y = rgdp[,2], data.ydate = rgdp[,1],
  data.x = ads[,2], data.xdate = ads[,1],
  est.start = as.Date(est.start),
  est.end = as.Date(est.end), horizon = 22, x.lag = 66,
  y.lag = 1, polynomial = "nealmon",
  method = "fixed", disp.flag = TRUE,
  num.evals = 10000, num.coef = 10)
```

Compute U-MIDAS nowcasts

```
ads.umidas.obj.fixed.now <- midas.obj.fixed.for <- midas.adl(data.y = rgdp[,2],
  data.ydate = rgdp[,1], data.x = ads[,2],
  data.xdate = ads[,1], est.start = as.Date(est.start),
  est.end = as.Date(est.end),
  horizon = 22, x.lag = 66, y.lag = 1,
  polynomial = "umidas", method = "fixed",
  disp.flag = TRUE)
```

Compare forecasts

```
ads.midas.obj.fixed.for$pred.obj$rmse
```

```
[1] 0.3096144
```

```
ads.umidas.obj.fixed.for$pred.obj$rmse
```

```
[1] 0.6149786
```

Compare nowcasts

```
ads.midas.obj.fixed.now$pred.obj$rmse
```

```
[1] 0.3486505
```

```
ads.umidas.obj.fixed.now$pred.obj$rmse
```

```
[1] 0.7206917
```

Take away

MIDAS regression compares well with alternative methods, i.e. time-averaged data or U-MIDAS specifications.

Even when sampling frequency ratio is small, e.g. quarterly/monthly, more often than not MIDAS regression with tightly parameterized polynomial yields better quality forecasts, mainly through the bias-variance trade-off argument.

NLS estimation is cumbersome, sometimes gives local optimum parameter estimates due to nonlinearities of the objective function. In the case of ARDL-MIDAS, these can be overcome by using profiling approach, see Ghysels and Qian (2019).